AI INFRASTRUCTURE LLM SERVING VLLM CUDA RAY SERVE

AI Inference Infrastructure Engineer Roadmap [2026]

Master vLLM, TensorRT-LLM, KV caching, and Ray Serve to secure the highest-paying AI Inference Infrastructure Engineering role of 2026.

TI techinterview.net
June 2026

30 min read

In 2023, the tech world was obsessed with *training* Large Language Models. Companies poured billions of dollars into massive H100 GPU clusters just to compute the weights of foundation models. But in 2026, the paradigm has shifted. Foundation models are ubiquitous, and the competitive moat is no longer just about building the smartest model—it's about serving it at scale, with ultra-low latency, without going bankrupt.

Welcome to the era of the **AI Inference Infrastructure Engineer**.

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

MANAGE PREFERENCES REJECT NON-ESSENTIAL ACCEPT ALL

the brutal realities of LLM serving, the memory wall, PagedAttention, quantization strategies, and the dominant orchestration frameworks.

---

1. The Cruel Reality: The Memory Wall

If you come from traditional microservices (Java, Go, Node.js), your intuition for performance tuning is fundamentally wrong when it comes to AI inference. In traditional backend engineering, the bottleneck is almost always CPU or network I/O. In LLM inference, the CPU is practically irrelevant, and the GPU's compute cores (TFLOPS) are often sitting idle.

The ultimate enemy of LLM inference is **Memory Bandwidth**. To understand how to design around this constraint, read our **GPU Memory Architecture System Design Guide →**.

Large Language Models are memory-bound. Generating a single token requires the GPU to read the *entire* model's weights from High Bandwidth Memory (HBM) into the GPU registers. If you have a 70 Billion parameter model running at 16-bit precision (FP16), the model weights alone consume 140GB of VRAM. A single Nvidia A100 GPU only has 80GB of VRAM. This means you literally cannot even fit the model onto a single chip; you must shard it across multiple GPUs using Tensor Parallelism. (Alternatively, new hardware like the AMD MI350X solves this with 288GB of HBM3E memory; for a deep dive, see our **AMD ROCm 10 & MI350X Interview Guide →**). Google

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

2. The KV Cache and PagedAttention

To understand why LLM inference is hard, you must understand the Key-Value (KV) Cache.

LLMs generate text autoregressively (one token at a time). When generating the 100th token, the model needs the mathematical context of the previous 99 tokens. Recomputing the attention mechanism for all 99 tokens on every single step is mathematically redundant and brutally slow.

Instead, inference engines cache the key and value tensors of past tokens. This is the **KV Cache.**

As a model generates a long response, the KV Cache grows dynamically. If you have 100 users querying your model simultaneously, each user has their own massive, unpredictable KV Cache. Historically, engines allocated contiguous blocks of GPU memory for the *maximum possible* sequence length. If a user asked a short question, 95% of that allocated memory was wasted. Because GPU VRAM is the most expensive real estate on earth, this fragmentation crippled batch sizes.

The vLLM Revolution: PagedAttention

In late 2023, researchers at UC Berkeley introduced **vLLM** and a technique called **PagedAttention.** Inspired by classic operating system virtual memory management, PagedAttention divides the KV Cache into fixed-size blocks (pages). To see how cutting-edge architectures completely overhaul

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

Mastering vLLM is step one on this roadmap. If you are preparing for an interview at an AI infrastructure company, expect grueling whiteboarding sessions on PagedAttention logic. To master these mechanics, tackle our **vLLM Inference Engine Interview Questions →**. To test broader distributed systems knowledge, review our **AI Infrastructure & MLOps Interview Question Bank →**.

3. The Engine Wars: vLLM vs. TensorRT-LLM vs. TGI

In 2026, you cannot just deploy a PyTorch script to production. You must use a highly optimized inference server. The three titans of the industry are vLLM, TensorRT-LLM, and TGI. An Inference Engineer must know precisely when to use each.

Feature / Engine	vLLM (UC Berkeley / Open Source)	TensorRT-LLM (Nvidia)	Text Generation Inference (TGI) (Hugging Face)
Primary Advantage	Industry-standard PagedAttention, easiest to set up, massive	Absolute peak performance on Nvidia hardware,	Excellent integration with the Hugging Face ecosystem,

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

			(Hugging Face)
Quantization	AWQ, GPTQ, FP8	FP8, INT8 SmoothQuant, AWQ	BitsAndBytes, GPTQ, AWQ
Best Use Case	Startups, rapidly iterating on new models, generalized API endpoints.	Massive enterprise scale where squeezing out 5% more throughput saves millions.	Managed deployments closely tied to the Hugging Face hub.

The Nvidia Moat: TensorRT-LLM

While vLLM is the darling of the open-source community, **TensorRT-LLM** is what the giants (OpenAI, Anthropic, Meta) use when they need to squeeze every last drop of blood from an H100 cluster.

TensorRT-LLM isn't just an engine; it's a compiler. It takes your PyTorch model and fuses the operations (Kernel Fusion) specifically for the architecture of your target GPU. It utilizes highly proprietary Nvidia libraries like cuBLAS and cuDNN. The learning curve is notoriously steep because you must pre-compile engines ahead of time, but the throughput gains at high concurrency are unmatched.

TensorRT-LLM is heavily integrated with Triton. To master these bare-metal C++ deployments,

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

4. Model Quantization: Compressing the Brain

If memory bandwidth is the enemy, quantization is your ultimate weapon.

Quantization is the process of reducing the mathematical precision of the model's weights and activations. A standard model uses 16-bit floating-point numbers (FP16). By compressing these to 8-bit integers (INT8) or even 4-bit integers (INT4), you drastically reduce the memory footprint.

- **FP16 (16-bit):** High accuracy, terrible memory usage.
- **INT8 (8-bit):** The modern sweet spot. Half the memory, almost imperceptible loss in quality.
- **INT4 / AWQ (4-bit):** Extreme compression. A 70B model fits on a single 80GB GPU. Great for edge devices or consumer hardware, but can suffer from "hallucination creep."

The Mechanics of Quantization (AWQ vs. GPTQ)

Not all quantization is created equal. Naively rounding numbers down destroys the intelligence of the model.

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

5. Orchestration: Scaling to a Million Users

So, you have vLLM running a quantized model on a single node. Congratulations. Now, how do you handle 10,000 concurrent requests across a fleet of 500 GPUs distributed globally?

Traditional load balancers (like NGINX or HAProxy) fail spectacularly at LLM inference. Why? Because LLM requests are wildly heterogeneous.

User A asks the model to output the word "Hello" (takes 10 milliseconds). User B asks the model to write a 4,000-word essay (takes 30 seconds). A traditional round-robin load balancer will eventually route a new request to a node that is currently blocked writing an essay, causing a massive latency spike for the user, while other nodes sit idle.

The Solution: Ray Serve

In 2026, **Ray** is the undisputed king of distributed AI execution. Originally built by researchers at UC Berkeley, Ray handles the chaotic statefulness of AI workloads perfectly.

Ray Serve sits above your inference engines (like vLLM). It acts as an intelligent, asynchronous

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

requests into the batch the exact millisecond an older request finishes generating its final token. To master the system design patterns required to build these resilient, auto-scaling inference orchestration layers, study our **Serverless GPU System Design Interview Guide →**.

Understanding distributed systems at this scale is arguably the hardest part of the interview process. Before applying, you must master these multi-node orchestration concepts in our **Ray Distributed AI Interview Questions →**. To see how emerging deterministic hardware architectures bypass the need for complex runtime batching and orchestration altogether, study our **Groq Inference Engineer Interview Guide →**.

---

6. The Physical Reality: Networking and Hardware

Software can only take you so far. When you are spanning a massive language model across 8 GPUs in a single physical node (Tensor Parallelism), the bottleneck moves from software to physical interconnects.

PCIe Gen 5 is too slow. If GPU 1 needs to pass a 10GB tensor to GPU 2 via the motherboard's PCIe bus, the latency destroys the generation speed. This is why Inference Engineers must understand hardware topology.

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

(Remote Direct Memory Access). To see how companies like xAI scale these RDMA fabrics to 100k GPUs using RoCEv2 instead of InfiniBand, check out **xAI Interview Questions →**. For a deep dive into the RDMA data paths and hardware-enforced flow control that make this possible, read our **InfiniBand vs RoCEv2 Network Guide →**.

AI Cluster Networking: Understanding the strict boundary between NVLink and Ethernet is critical for deploying multi-node inference. Master the differences between **Scale-Up vs. Scale-Out networking** in our **RDMA & RoCEv2 AI Cluster Network Design Guide →**.
To bridge the gap between traditional DevOps reliability and the harsh realities of physical GPU infrastructure, SREs should transition their mental models by studying our **SRE & Cloud Infrastructure Interview Guide →**. If you are fascinated by how code interacts with raw silicon and real-world machinery, this discipline closely mirrors the challenges found in the **Physical AI Engineer Roadmap →**.

---

7. The 2026 Blueprint: How to Get Hired

The role of an AI Inference Infrastructure Engineer sits perfectly at the intersection of three brutal disciplines: Distributed Systems, Low-Level Systems Programming (C++/CUDA), and Machine Learning.

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

Month 2: The Math of Compression Learn the math behind quantization. Use an open-source library like AutoAWQ to compress a 16-bit model down to 4-bit. Serve it. Measure the memory footprint difference. Monitor the VRAM usage via 'nvidia-smi'.

Month 3: Cluster Scale Learn Ray. Deploy a multi-node Ray cluster. Use Ray Serve to act as an asynchronous frontend for multiple vLLM instances. Write a load-testing script using Python's 'asyncio' that blasts your cluster with heterogeneous requests (some asking for 10 tokens, some asking for 2,000 tokens). Watch how Ray handles continuous batching.

When you sit in the interview room, you won't just be reciting algorithms. You will speak from the battle-hardened experience of having actually wrestled with the Memory Wall. In 2026, that is the exact talent the industry is desperately willing to pay for.

Frequently Asked Questions

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

What is the difference between TensorRT-LLM and vLLM?

Why use Ray Serve for LLM orchestration?

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

DEEPEN YOUR PREP: MORE FROM THE BLOG

SPECULATIVE DECODING · LLM INFERENCE

Speculative Decoding: LLM Inference Guide [2026]

30 MIN READ · READ ARTICLE >

VLLM · PAGEDATTENTION

vLLM & PagedAttention: KV Cache Interview Guide [2026]

30 MIN READ · READ ARTICLE >

FLASHATTENTION-3 · TRITON

FlashAttention-3 & GPU Kernel Interview Guide (2026)

32 MIN READ · READ ARTICLE >

RDMA · ROCEV2

RDMA & RoCEv2: AI Cluster Network Design [2026]

32 MIN READ · READ ARTICLE >

SYSTEM DESIGN · MACHINE LEARNING

LoRA & PEFT Serving System Design: S-LoRA [2026]

36 MIN READ · READ ARTICLE >

MACHINE LEARNING · SYSTEM DESIGN

Tensor vs Pipeline Parallelism: 3D AI Architecture [2026]

30 MIN READ · READ ARTICLE >

NETWORKING · AI INFRASTRUCTURE

InfiniBand vs RoCEv2: CoreWeave AI Cluster Design [2026]

SYSTEM DESIGN · FINE-TUNING

LLM Fine-Tuning System Design Interview Guide: [2026]

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.

- DeepSeek MoE Architecture Interview Questions
- NVIDIA NIMs Interview Questions

Master Your Next Interview

Test your skills with AI-powered mock interviews tailored to your role.

START PRACTICING FREE

NO CREDIT CARD REQUIRED • INSTANT AI FEEDBACK

We value your privacy

We use cookies to improve your experience, secure our services, and analyze website traffic. By clicking "Accept All", you agree to our use of cookies. You can change your settings anytime.